

# General

- Trade-offs exist between **consistency, availability, and performance**.
- **Indexing strategies** significantly impact query performance.
- **NoSQL databases** are great for unstructured data but may sacrifice strong consistency.
- **Graph databases** are optimized for relationships.

## Large Scale Storage & Retrieval: Foundations

- **Foundations of Large-Scale Storage & Retrieval**

- Introduction to Information Storage**

- Information retrieval (IR) is the process of **storing, organizing, and retrieving data efficiently**.
    - IR systems range from **search engines (Google)** to **databases storing medical records**.
    - Key Challenges:
    - **Scalability**: Handling petabytes of data.
    - **Efficiency**: Fast retrieval times (e.g., Google in milliseconds).
    - **Consistency & Availability**: Ensuring fault tolerance and reliability.

- **Searching Strategies**

- **Linear Search**: Simple but inefficient ( **$O(n)$**  time complexity).
    - **Binary Search**: Only works on sorted data ( **$O(\log n)$**  complexity).
    - **Indexing**: Key to making search efficient; uses external data structures like **B-Trees and Hash Indexes**.

- **Indexing Techniques**

- **Hash Indexing**: Best for unique lookups ( $O(1)$  average case).
    - **B-Trees & B+ Trees**: Used in databases, optimized for disk access.
    - **Inverted Index**: Used in search engines, maps words to documents.

Record – A collection of values for attributes of a single entity instance; a row of a table

Collection – a set of records of the same entity type; a table

- Trivially, stored in some sequential order like a list

Search Key – A value for an attribute from the entity type

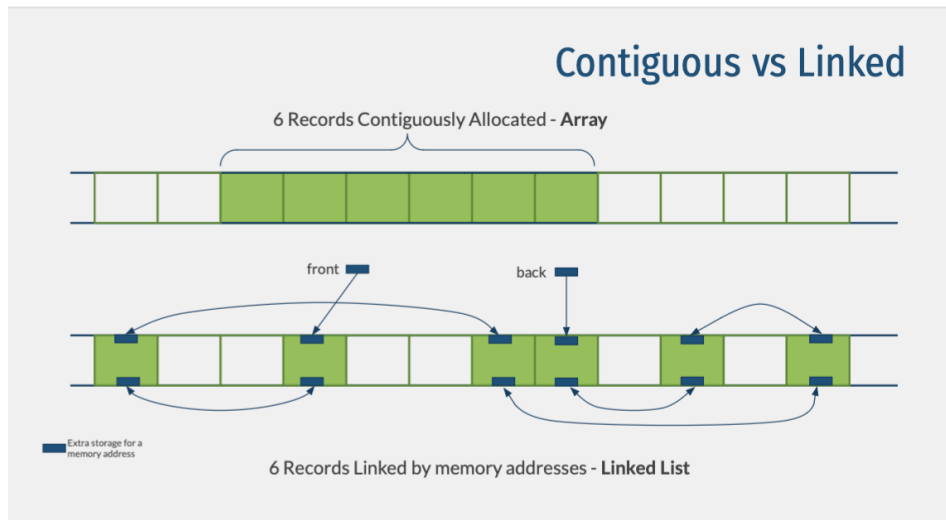
- Could be  $\geq 1$  attribute

$n \times x$  bytes of memory = If each record takes up  $x$  bytes of memory, then for  $n$  records we need

Contiguously Allocated List = All  $n \times x$  bytes are allocated as single "chunk" of memory

Linked List

- Each record needs  $x$  bytes + additional space for 1 or 2 memory addresses
- Individual records are linked together in a type of chain using memory addresses



Arrays are faster for random access but inserting at the end is slow

Linked Lists are faster for inserting anywhere in the list but random access is slow

Linear Search

Best Case = 1

Worst Case


**Binary Search**

Best Case = 1

Worst Case =  $\log_2 n$

*Read through this code it will explain what binary search is/does:*

```
def binary_search(arr, target)
    left, right = 0, len(arr) - 1
    while left <= right:
        mid = (left + right) // 2
        if arr[mid] == target:
            return mid
        elif arr[mid] < target:
            left = mid + 1
        else:
            right = mid - 1
    return -1
```



Binary Search Tree – a binary tree where every node in left subtree is  $<$  parent; every node in right is  $>$  parent

AVL Tree – Approx balanced binary search tree

- Maintains balance faster
- AVL Property  $|h(LST) - h(RST)| < \text{or} = 1$
- Self-balancing

Inserting Imbalance

- LL

- LR
- RL
- RR

## Trees

### . B-Trees and Indexing

#### Why B-Trees?

- Used in **database indexing** and **filesystems** (e.g., NTFS, PostgreSQL indexes).
- Keeps tree height small, reducing disk reads.

#### B+ Trees vs. B-Trees

- **B-Trees**: Store keys and values in internal and leaf nodes.
- **B+ Trees**: Store values only in leaf nodes, making range queries more efficient.

### BST Properties

- Left subtree values < Root < Right subtree values.
- Search time complexity:  **$O(\log n)$**  (balanced),  **$O(n)$**  (unbalanced).

### AVL Trees

- Self-balancing BST; ensures  **$O(\log n)$**  search time.
- Rotations used for rebalancing:
  - **LL Rotation** (Single Right Rotation)
  - **RR Rotation** (Single Left Rotation)
  - **LR Rotation** (Left-Right Rotation)
  - **RL Rotation** (Right-Left Rotation)

## SQL vs No SQL

### SQL

Structured Query Language

- Relational: easy querying on relationships on data between multiple datas
- ACID compliant: Atomicity, Consistency, Isolation, Durability

Not good for storing and querying unstructured data where format is unknown

Difficult to scale horizontally

### NoSQL

- Document, keyvalue graph, wide column stores
- More flexible and simpler to set up
- Can shard data across different data stores allowing for distributed databases, making horizontal-scaling much easier
- Very large data can be stored

Designed for distributed purposes

Optimistic Concurrency

- Assumes conflicts are unlikely to occur
- Add last updated timestamp and version number columns to every table that can be modified independently of transactions.

## Relational Databases vs. NoSQL

### ACID vs. BASE

- **ACID:** Ensures strong consistency.
- **Atomicity** - Transactions are all-or-nothing.
- **Consistency** - Data remains valid.
- **Isolation** - Transactions do not interfere.
- **Durability** - Once committed, stays committed.
- **BASE:** Prioritizes availability over consistency.
- **Basically Available, Soft-state, Eventually Consistent.**

### SQL vs. NoSQL

Feature	SQL (Relational)	NoSQL (Non-relational)
<b>Schema</b>	Fixed, structured	Flexible, dynamic
<b>Scalability</b>	Vertical (scale-up)	Horizontal (scale-out)
<b>Consistency</b>	Strong ACID Compliance	Eventual Consistency
<b>Examples</b>	MySQL, PostgreSQL	MongoDB, Cassandra, Neo4j

## NoSQL

- **CAP Theorem**
  - **Consistency:** All nodes see the same data simultaneously.
  - **Availability:** Every request gets a response (even if stale).
  - **Partition Tolerance:** System functions despite network failures.
- **NoSQL Categories**

Type	Example	Use Case
Key-Value Stores	Redis, DynamoDB	Caching, real-time analytics
Document Stores	MongoDB, CouchDB	JSON-like flexible data models

Column Stores	Apache Cassandra	Time-series & analytics
Graph Databases	Neo4j	Relationship-heavy data

## Graph Data Models & Neo4j

### Graph Theory Basics

- **Nodes** = Entities (e.g., users, products).
- **Edges** = Relationships between nodes.
- **Traversal Algorithms**: BFS, DFS, Dijkstra's Shortest Path.

### Neo4j & Cypher

- Neo4j is an **ACID-compliant** graph database.
- **Cypher Query Examples**:  
MATCH (a)-[:FRIEND]->(b) RETURN a, b;

## Data Replication & Distributed Systems

- Improves **fault tolerance** (if a node crashes, others still serve data).
- Enables **load balancing** (distribute requests across multiple servers).

### Replication Strategies

- **Master-Slave Replication**: One node handles writes; others replicate for reads.
- **Multi-Master Replication**: Multiple nodes accept writes; requires conflict resolution.
- **Eventual Consistency**: Updates propagate over time (**used in NoSQL systems** like DynamoDB)