

Binary Search Trees (BST)

- A BST is a binary tree where:
 - Left subtree nodes have keys less than the parent.
 - Right subtree nodes have keys greater than the parent.
- Keys are distinct.
- Each node has pointers to parent, left child, and right child.

Traversal Methods

- **Inorder:** Left subtree → Node → Right subtree (Outputs sorted order)
- **Preorder:** Node → Left subtree → Right subtree
- **Postorder:** Left subtree → Right subtree → Node

BST Operations

1. **Searching**
 - Compare the target key with the node's key and traverse left or right accordingly.
 - Time Complexity: $O(h)$, where h is the tree height.
2. **Insertion**
 - Search for the correct position and insert while maintaining BST property.
3. **Deletion**
 - Three cases:
 - Node has no children → Remove it.
 - Node has one child → Replace it with the child.
 - Node has two children → Replace it with its in-order successor.

BST Performance

- Worst-case height: $O(n)$ (degenerated tree).
- Average height for random BST: $O(\log n)$.

Balanced Trees

BSTs can become unbalanced, degrading search efficiency. Self-balancing trees maintain $O(\log n)$ height.

AVL Trees

- Self-balancing BST with height-balance property:
 - For every node, height difference between left and right subtrees ≤ 1 .
- Rotations:
 - **Single Rotation (LL, RR):** Simple rotation when imbalance occurs on one side.

- **Double Rotation (LR, RL):** Two rotations required when imbalance occurs across different sides.
- Time Complexity:
 - Search: $O(\log n)$
 - Insert/Delete: $O(\log n)$ (may require rebalancing)

B-Trees and B+ Trees

- Designed for disk-based storage and indexing.
- Generalization of BST where each node can have multiple children.
- Each node has up to **m** children (order **m** tree).
- Ensures that all leaf nodes are at the same depth.

B-Trees

- Internal nodes store keys and pointers to child nodes.
- Leaf nodes contain data records.
- Insertion and deletion maintain balance through **node splitting and merging**.

B+ Trees

- Variation of B-Trees where:
 - **Internal nodes only store keys** (not data).
 - **Leaf nodes store all data records** and are linked together.
 - Enables fast range queries and sequential access.

B+ Tree Insertion Process

1. Find the correct leaf node.
2. Insert the key.
3. If the leaf is full, **split** it and push the middle key to the parent.
4. If the parent is also full, **split** it and repeat up to the root.
5. If the root splits, a new level is created, making the tree deeper.

Conclusion

- **BSTs** are simple but can become unbalanced.
- **AVL Trees** keep balance at the cost of rotations.
- **B-Trees and B+ Trees** optimize disk-based searches, with **B+ Trees** excelling in range queries.
- Choice of structure depends on access pattern and data size.

Database Systems: Moving Beyond Relational Models & NoSQL

1. Relational Database Model Benefits

- **Standard model and language:** Widespread SQL adoption
- **ACID compliance:** Atomicity, Consistency, Isolation, Durability
- **Strengths:** Works well with structured data and can handle large data volumes
- **Established ecosystem:** Well-understood with extensive tooling and expertise

Performance Optimization Techniques

- Indexing
- Direct storage control
- Column vs. row oriented storage
- Query optimization
- Caching/prefetching
- Materialized views
- Precompiled stored procedures
- Data replication and partitioning

2. Transaction Processing & ACID Properties

Transaction Fundamentals

- **Definition:** Sequence of CRUD operations performed as a single logical unit
- **Outcomes:** Either fully succeeds (COMMIT) or entirely fails (ROLLBACK)
- **Benefits:** Data integrity, error recovery, concurrency control, reliable storage

ACID Properties in Detail

- **Atomicity:** Transaction is an indivisible unit - fully executed or not at all
- **Consistency:** Database moves from one consistent state to another
- **Isolation:** Concurrent transactions don't interfere with each other
 - Potential issues: Dirty reads, non-repeatable reads, phantom reads
- **Durability:** Committed changes persist even during system failures

Isolation Issues

1. **Dirty Read:** Transaction reads uncommitted data from another transaction

2. **Non-repeatable Read:** Two queries in same transaction get different values due to another committed transaction
3. **Phantom Reads:** Records appear/disappear during a transaction due to another transaction's inserts/deletes

3. Limitations of Relational Databases

- Schema evolution challenges
- Full ACID compliance not always necessary
- Expensive join operations
- Limited handling of semi-structured/unstructured data (JSON, XML)
- Horizontal scaling difficulties
- Performance limitations for real-time/low-latency systems

4. Scaling Approaches

- **Vertical Scaling (Up):** Using more powerful systems
 - Easier implementation, no architecture changes
 - Practical and financial limitations
- **Horizontal Scaling (Out):** Distributed computing model
 - Modern systems making this more feasible
 - Requires handling distributed data storage

5. Distributed Data Systems

Characteristics

- Computers operate concurrently
- Independent failure modes
- No shared global clock

Data Storage Approaches

- Data stored on multiple nodes, typically replicated
- Can be relational or non-relational
- Network partitioning is inevitable
- System needs partition tolerance

6. CAP Theorem

Core Principle

It's impossible for a distributed data store to simultaneously provide more than two of:

- **Consistency:** Every read receives most recent write or error
- **Availability:** Every request receives a non-error response
- **Partition Tolerance:** System operates despite network issues

Practical Interpretations

- **C+A:** Latest data, guaranteed response, but vulnerable to network partitions
- **C+P:** Always latest data or no response, handles partitions
- **A+P:** Always responds but may not have latest data, handles partitions

Reality

- Not about always giving up one property
- More about trade-offs when partitions occur

7. NoSQL Databases

Concurrency Models

- **Pessimistic Concurrency (ACID):** Assumes conflicts will occur, uses locks
- **Optimistic Concurrency:** No locks, assumes conflicts are rare
 - Uses timestamps and version numbers
 - Good for read-heavy systems
 - Less efficient for high-conflict systems

BASE - Alternative to ACID for Distributed Systems

- **Basically Available:** System works most of the time
- **Soft State:** State can change without input
- **Eventual Consistency:** System will eventually reach consistency

8. Key-Value Databases

Design Principles

- **Simplicity:** Extremely simple data model compared to RDBMS
- **Speed:** Often in-memory, O(1) retrieval operations
- **Scalability:** Easy horizontal scaling with eventual consistency

Use Cases

- **Data Science:** Experiment results storage, feature storage, model monitoring
- **Software Engineering:** Session information, user profiles, shopping carts, caching

9. Redis - Key-Value Database

- Open-source, in-memory database
- High performance (>100,000 SET ops/second)
- Supports data durability through snapshots or append-only journals
- Simple lookup by key only (no secondary indexes)

Data Types

- **Strings:** Basic text/binary data
- **Lists:** Linked lists for queues/stacks
- **Sets:** Unique unordered string collections
- **Hashes:** Field-value collections
- **Sorted Sets:** Sets with scores for ordering
- **Specialized Types:** Geospatial data, JSON, etc.

Common Operations

- Basic CRUD: SET, GET, EXISTS, DEL
- Atomic counters: INCR, INCRBY, DECR, DECRBY
- List operations: LPUSH, RPUSH, LPOP, RPOP, LLEN, LRANGE
- Set operations: SADD, SCARD, SISMEMBER, SINTER, SDIFF
- Hash operations: HSET, HGET, HGETALL